Constraint-Emergence Ontology — v2.0 (Unified Cut)

Dimitar Popov

July 13, 2026

Contents

1 Reading Contract

1.1 What kind of document this is

Read this as a constitution. It is written to be loaded into a language model's context window, where its job is to *constrain* — to shape what the model explores the way a set of axioms shapes what a geometer can prove. Truth of the kind a measurement has is the wrong success test for a document with that job. Three intents stack here, and the third governs. It is philosophy, in the old sense that made chemistry thinkable before there were atoms to weigh and made natural selection arguable before there were genes to find; its falsifiability markers are declared exposure, and grading its pre-formal claims by the standards of a lab misreads them. It is systematic construction in the continental line of Whitehead, Simondon, and Deleuze, where a metaphysics is judged by whether it coheres and generates. And decisively, it is an *experimental axiomatic ontology*: a compact frame whose success question is **does this axiom set, loaded as context, make exploration more coherent within the set, more transferable across domains, and more novel than exploration without it?** The first two intents ask "is this true physics?" and "does the system hang together?"; the third asks only this, and the harness that closes this document settles it. Argument cannot.

1.2 The five-status discipline and the cut law

Because the axioms here are a constraint system, ambiguity is a functional defect: an ambiguous constraint produces drifting exploration. So two disciplines run throughout. The first is that every substantive claim carries exactly one of five epistemic statuses — marked inline in the prose as the claim is made — and they are the whole vocabulary: an **axiom** is a generative first principle you reason *from*; to reject one is to exit the framework, which is why an axiom owes no proof and admits no refutation; a **theorem-within-set** is earned by derivation from the axioms and can be challenged only on that derivation; a **conjecture** is an open claim inside the set, plausible but undecided, whose consequences we explore both ways; an **empirical-exposure-point** is a claim the framework stakes against outside data, the place where it can actually lose; and an **interpretive-stance** is a way of seeing compatible with existing science but carrying no independent test, to be used strictly as a lens; asserting one as discovered fact violates its status. The second discipline is that the framework is *cut* like a versioned ledger: each revision is a numbered cut carrying a digest of what changed and why, and a cut does not close until the evaluation harness has run on it. This is v2.0, a draft cut that restores the full cumulative traversal of v1 inside the typed, status-disciplined frame developed during revision — the vocabulary and axioms are ratified for use, the harness is specified but not yet executed, and that gap is this cut's declared reason for remaining open.

1.3 Why the traversal must be watched

One reading instruction governs the order of what follows. The framework's central claim is that one structure recurs across physics, biology, language models, and software. That claim is not stated up front and then illustrated; it is *precipitated*. The physics traversal comes first and does one job: it manufactures the vocabulary — every typed term gets defined against a concrete physical object before it is asked to travel. Then the rivals are met, then the hardest tests are declared, and only then does the same structure get shown crossing into biology, language models, and software. The reader is meant to watch the shape appear three times over before the formal category that names it is written down. Watching it precipitate is the argument; naming it early would be assertion.

2 The puzzle we walk into

2.1 The hidden Platonic premise

Start with a drift in how physics talks about itself, because the whole framework is a way out of the corner that drift paints. Since Galileo announced that nature was written in mathematics, one strand of physics has slid toward treating the writing as the thing written. Newton's laws hardened from a description of how planets move into the reason they move; the twentieth century abstracted further — Hilbert spaces, gauge symmetries, renormalization — until Tegmark could say plainly that reality *is* a mathematical structure. Many-Worlds is the terminus of that drift. If the formalism is ontologically primary, then everything the formalism contains must exist; the wavefunction carries superposition terms, so every branch is real, and the actuality you happen to observe is a parochial illusion. This is coherent. It may even be true. But notice what forces it: a hidden premise smuggled in underneath — that mathematical existence confers ontological existence, that what can be written down, exists. The physics itself, which every interpretation shares, forces nothing of the kind.

2.2 Computation instead of enumeration

The premise underneath everything that follows: reality is computation over constraints, and the rival programs are attempts to replace the computation with its completed value. String theory and Many-Worlds each try, in different registers, to construct every universal state — to hold the whole state space as a finished mathematical object. We take the opposite wager: the state space is never constructed, only traversed. What exists is what has been computed — each admissible step actualizing one configuration against the constraint network — and the catalogue of what might have been computed exists only as the constraints' silent shape. Enumeration is a description of computation; it is never the computation.

2.3 Computational presentism

That wager can be derived, not merely asserted, and the derivation is worth walking because it fixes what "exists" is allowed to mean for the rest of the document. The constraint network evolves in discrete steps. Each step produces the next state from the current one. Take this as the derivation of **computational presentism** — a theorem within the set once the discrete-step axiom is granted: the current state is the only state that exists, because the prior state was consumed in producing it and the next state has not yet been computed. The past does not exist; only the results of prior computation persist as the present configuration. The future does not exist; it has not been computed, and the network's next state is undetermined until the next step occurs. This is what Smolin reaches for when he says time is real, sharpened into a claim about existence rather than about time: to exist is to be the output of the constraint network's forward evolution, and a completed state space — every branch, every path, held at once — is a description mistaken for that output. The rival programs enumerate; the network computes; and enumeration is never the computation.

2.4 The map remains a map

We decline the Platonic premise, and the whole document unfolds from the declining. Take this next move as an interpretive stance — a lens compatible with all the physics, carrying no independent test of its own: our laws of physics are descriptions of physics, maps of a territory. However good the map, it remains a map; the planets do not consult F=ma before moving. On this reading the wavefunction describes a *potentiality structure* — Aristotle's sense in which an acorn genuinely holds the potential of an oak without there being a second, ghostly oak standing beside it, and equally Deleuze's *virtual*, which is real without being actual. Superposition is real potentiality. Collapse is real actualization of that potentiality into one configuration. The branches that failed to actualize remain what they always were: possibilities. Many-Worlds reifies a description into a substrate; it mistakes the enumeration for the execution. We do the opposite subtraction. This is the one anchoring stance of the framework, and it is worth being clear about its cost and its scope: it buys nothing empirical, generates no unique number, and every physics-mechanism elaboration that hangs from it — variable-c, the Higgs as order parameter, gravity as constraint density, the proton-electron ratio — earns reputational exposure without earning predictions, and is carried in this document as declared interpretive framing, never as result. What survives as load-bearing is the single move: descriptions describe, and we should stop mistaking a beautiful compression for the thing it compresses. ("Platonic" and "Aristotelian" here are shorthand for two poles of a perennial tension — ideas *are* reality versus ideas *describe* it — and carry no exegetical weight.)

If the equations merely describe the substrate, the obvious question is what the substrate is. The physics traversal that follows answers it, and in answering it manufactures every word the rest of the framework will need.

3 The physics traversal, where the vocabulary is made

3.1 From substance to admissible transformation

The honest answer to "what is reality made of?" is that the question presupposes stuff, and stuff keeps dissolving under inspection — particles into fields, fields into excitations, excitations into whatever the next formalism says. So invert the question, from what things *are* to what transformations are *admissible*: what can follow from what, what compositions hold, what is forbidden. This is a shift from objects to morphisms, the categorical instinct — attend to what things do and what structure is preserved under mapping, and the question of ultimate composition dissolves with its subject. Make that shift consistently and a candidate substrate appears that is free of stuff altogether.

3.2 Axiom C and the one meaning of constraint

This we take as axiom — Axiom C, the constraint-network ontology. The substrate is an evolving global constraint network: a self-consistent system of admissible and forbidden transitions. What we call fields, particles, and spacetime are its stable patterns and projections. The invariants of reality live in the structure of admissible transformation — structural realism relocated from objects to dynamics; the substance transformed carries none of them. This is Bohm's move made precise. The constraint network is his implicate order, the enfolded substrate; the fields and particles we handle are the explicate order unfolded from it; and the framework's distinction between a stable pattern and a projection is his enfolding read as substrate-versus-emergent-manifold. Bohm reached a deterministic, non-local substrate decades before Bell's theorem was widely understood, and described it through the hologram and the ink-drop; this framework's job is to characterise the same substrate structurally. Rejecting Axiom C steps outside the framework. No refutation occurs at that step.

Fix, then, the one word the whole system turns on, because letting it mean six things on a page was the central functional bug of v1.3. A **constraint** — the bedrock term, taken as axiom — is a condition that determines which transformations are admissible. Admissibility exhausts it; a rule dictating what happens next is a different kind of thing. That is its *only* meaning here. Where the earlier work let "constraint" also be a potential well, a test suite, a prompt, and a metric, those are now named as distinct images of admissibility, and the bare word is reserved for the axiom's sense. The strictness is functional: in a constitution the vocabulary *is* the constraint system the model explores within, so a word meaning six things is a manifold with six overlapping metrics, and exploration inside it drifts. Constraint also *generates*, beyond what it permits. Reality is the remainder of the forbidden, the structure that survives when the inadmissible is excluded.

3.3 Boundaries, hydrogen, and Markov objects

Now the first typed object, taught against something you can hold. A pattern persists only if it has an inside that goes on being itself somewhat regardless of the churn outside, which is to say it has a *boundary* across which the outside is screened off. The hydrogen atom is the cleanest instance. Its electron obeys a wave equation in a potential well; the admissible solutions are normal modes, harmonics on a drum, with fixed spatial structure and quantised energy. An atom is a quantum resonant cavity, and its potential well is a **boundary** — axiomatic, one of the framework's typed primitives — the partition into interior, boundary, and exterior that screens the interior's dynamics from external perturbation. The atom is therefore a **Markov object**: a stable pattern whose interior is conditionally independent of the exterior given the boundary state, `P(interior | boundary, exterior) = P(interior | boundary)`. This screening property is Friston's Markov blanket stated substrate-neutrally, and this is the one place in the document it is written out in full; everywhere after, "Markov object" carries it by reference. "Markov object" is what a Markov blanket becomes once you stop assuming it only wraps living things.

Axiom M, Markov objects are the units of stable form. Wherever the constraint structure can sustain a boundary with the screening property, a stable pattern exists there. In physics these are standing waves, bound states, particles — eigenmodes of the constraint network, long-lived where the geometry permits. Fields, on this reading, are the constraint geometry itself: grades of mesh, sectors of the constraint manifold with their own density and topology, within which Markov objects precipitate where the mesh allows a stable pattern. The field is the mesh; the particle is the pattern the mesh permits. The *existence* of Markov objects wherever constraints permit is axiomatic; the *precision* of the screening in any real substrate is an empirical question, and that is why Axiom M carries an exposure marker — we cash it as a testable claim later, and it can lose.

3.4 Three processes hidden inside collapse

Collapse is the next term, and precision here repairs a conflation the earlier work carried. **Collapse is the resolution of admissible potentialities into one actual configuration under constraint over-determination.** In physics it reads as three distinguishable processes that a single word must not blur. First, objective reduction: when a superposition would require incompatible constraint-density configurations — what General Relativity describes as incompatible spacetime geometries — the network cannot sustain the inconsistency and resolves, because self-consistency is a hard constraint, not because anything looked; this is the constraint reading of Penrose's gravitational threshold. Second, decoherence: below that threshold, entanglement with a vast environment makes phase information unrecoverable in practice, which is over-determination by the environment rather than by self-consistency failure. Third, information projection: an embedded observer updates its description as alternatives become inaccessible, which is epistemic, the observer's compression updating rather than a physical event. The first is objective and physical, the second objective and statistical, the third epistemic. Take the physics reading of collapse as interpretive stance; it earns a stronger, derived status later in the compute regimes, and the status law lets the one word carry both without ever wearing both at once.

3.5 Motion as pattern propagation

Motion, in a substrate with no stuff, needs its own account, and the account defines *traversal* concretely before the term generalises. If particles are standing patterns, their motion is pattern propagation: nothing translates through a pre-existing space; the pattern's influence propagates through adjacent constraint regions, each node updating on its neighbours. Electricity in a wire is the teaching case. The electron drift velocity is around a tenth of a millimetre per second — the thing barely moves — while the signal propagates at some seventy percent of the speed of light, because what travels is the electromagnetic disturbance passed neighbour to neighbour. Sound in air, waves on water, phonons in a lattice: in every case the thing stays mostly put and the pattern moves. On this reading c is the constraint network's propagation rate, the maximum speed at which one node can influence its neighbours, appearing as the universal constant when projected onto emergent coordinates. This is an interpretive stance, compatible with field theory and offering no unique number; its value here is that it makes "traversal" mean something concrete before the term is asked to carry gradient descent and least action.

3.6 Constraint density and the nerf-ball picture

The nerf-ball picture gathers these into one image and names the generative principle behind them. Imagine a space packed with soft balls of many sizes — the constraints — jostling and reconfiguring. The balls are an analogy for a topology we do not see directly. Reality happens in the gaps between them: the standing waves live in the space the constraints leave, shaped by the boundaries the constraints impose. This is Deacon's absential causation rendered spatially — what matters is not the constraints themselves but what they *exclude*, and the patterns that form in the excluded remainder. A quantum of energy propagating through the churning topology conserves its volume — charge, energy-momentum, quantum numbers — while its shape is dictated by the gaps available at each instant. Where the balls pack densely, the gaps are smaller in coordinate terms, so the same standing wave spans fewer coordinate units and looks smaller from outside; where they are sparse, it spreads. Take the whole gravity-as-density picture that follows from this — dense constraint regions slowing local update rates, the same pattern contracting in external coordinates — as interpretive stance. It reproduces General Relativity's phenomena in constraint language and predicts no deviation from it. What the image earns is the vocabulary: constraint density as a *metric* on the manifold, and boundaries as what generates form by exclusion.

3.7 Stratified time and the computed present

Time, in this substrate, is stratified rather than singular. Each layer of the constraint hierarchy defines its own characteristic timescale — the rate at which patterns on that constraint plane update. Quantum fields update on one scale, atoms on another, chemistry on another, organisms and ecosystems and cognition each on their own; faster layers look frozen to slower observers, slower layers look like static background to faster processes. What we experience as time at any scale is the measure of change on that scale's emergent constraint plane. Take stratified time as interpretive stance: it reframes existing physics and makes no prediction against it. Its role in the vocabulary is to make "emergence" mean passage between constraint planes, not accumulation of complexity.

Beneath all the stratified times sits the derivation already given — the unit of change and computational presentism. The base layer evolves in discrete units; every emergent time is a derivative measure of change on some plane above it; the current state is the only state that exists; and the base unit's character is epistemically inaccessible from within emergence, so the framework takes no position on whether it is constant, periodic, or temporal in any sense we would recognise. This is the substrate reading of the presentism theorem, and it is what forbids the completed state space that the rival programs assume.

3.8 Recursion and the limits of internal derivation

One recursion result belongs here, because it is where the framework's computation-over-derivation commitment shows its teeth. Gödel, Turing, and Cantor are usually treated as three separate results in three separate domains. Take it as a theorem within the set that they are one result: any system complex enough to refer to its own operation cannot completely characterise itself from within. Gödel's formal system encodes statements about its own provability and generates truths it cannot prove; Turing's halting evaluator must simulate programs including itself and cannot exist; Cantor's enumeration must account for its own diagonal and stays forever incomplete. The common structure is recursion as unbounded self-generated state: each self-referential step creates new state the system must account for, and accounting for it requires another step, faster than the system can exhaust. This is not a limit on truth. The Gödel sentence is true, the halting question has an answer for any specific program, the diagonal number exists. The limit is on *derivation from within*. The resolution is always the same — compute from outside: the mathematician evaluating a Gödel sentence is a physical constraint-satisfying process running the evaluation from a level the formalism cannot access.

3.9 Penrose and computation from outside

That resolution answers Penrose directly, and the answer is load-bearing rather than incidental. Penrose argues that human mathematical understanding involves non-computable processes, localised in gravitational reduction events, on the grounds that Gödel shows no formal system reaches all mathematical truth while the mind seems to. The argument stands on mathematical Platonism — truths existing independently, requiring a non-algorithmic channel to reach. The framework rejects that premise. Mathematics is computation, a program running in physical brains; when a mathematician sees that a Gödel sentence is true, no Platonic realm is accessed, and the brain is simply a computational system not trapped inside the self-referential loop it is evaluating. Gödel shows that self-referential derivation has hard limits. He does not show that an external computation does. No non-computability is required. This connects to Wolfram's computational irreducibility as its positive face: some processes cannot be shortcut, and the only way to know their output is to run them, which is *why* formal systems cannot reach certain truths — the truth is the output of an irreducible computation. The constraint network sits on the computation side of this divide throughout. It does not derive; it enforces consistency globally and simultaneously. This is also why, later, a demand to "prove the ontology mathematically" for its *substrate* is a category error — it asks the framework to justify its most fundamental primitive in one of its own downstream projections, like demanding a cellular automaton justify Newtonian mechanics without leaving the lattice rules. The category error applies only to the substrate. Proving structural correspondences *between* domains is standard mathematics inside emergence, and the framework owes exactly that proof later.

3.10 Physics claims remain projection-specific

One law governs every physics reading made in this traversal, and stating it before the vocabulary travels keeps the stances from borrowing on each other's credit. Each interpretation is statused per projection. The constraint reading of collapse, c as the network's propagation rate, gravity as constraint density, stratified time, the Higgs as order parameter later on — each carries its own marker at the place it is made, each stands or falls on its own exposure, and none inherits standing from a neighbour that happens to share its vocabulary. A stance that held up would confer nothing on the stance beside it, and a stance that fell would take nothing down with it except itself. The vocabulary is one system; the physics interpretations projected through it are separately statused wagers.

The vocabulary is now made: constraint, boundary, Markov object, collapse, traversal, and the emergence that passes between planes. Every one was defined against a physical object before being asked to travel. Two things remain before the terms leave physics — the drive that populates the manifold, and the two tests where the physics reading is hardest.

4 The generative principle, traversal, and emergence

The tempting account of stable structure is that it is rare, improbable, the lucky product of fine-tuning. The framework refuses that, and the refusal is its metaphysical heart.

4.1 Axiom G: possibility populates the manifold

Axiom G, the generative principle. As soon as a stable configuration is *possible* within a constraint structure, it *will* emerge. Emergence is exhaustive: the constraint manifold fills its own possibility space. Read computationally, this is lazy evaluation at universal scale — a configuration is actualized when the traversal reaches it, and the possibility space is never materialized as a whole; it exists only as the constraints' silent shape. Particles, life, encodings, constructors, computation — each is what constraints do when they can, arrived free of coaxing or luck. An analytic reader will hear a tautology here, possible and stable and will-emerge chasing each other in a circle, and read as an empirical generalization it *is* vacuous. It stands here as a named metaphysical ultimate, in the exact sense of Whitehead's "creativity," the category of the ultimate that explains but is not itself explained. It is the same move as Simondon's drive of a metastable field toward individuation — individuation from a pre-individual field *is* Markov-object formation from a constraint network, articulated in 1958 — and the same move as Kauffman's insistence that systems saturate the adjacent possible. An axiom that grounds a system is supposed to be unfalsifiable in this way; testing it against data is a category error.

4.2 Axiom T: traversal follows the local slope

Axiom T, traversal is the one engine. All structured change is local preorder traversal on a constraint manifold: at each step, evaluate the local preorder — the preferred direction, the way certain states are more satisfied or lower in some potential than their neighbours — and move that way. Least action in physics, gradient descent in training, a model integrating its learned direction field in inference, a builder reducing the gap to a definition of done: instances of one operation, and the identity is asserted literally, because analogy would be the weaker claim and Axiom T makes the stronger one. The unit of change is the discrete step of this engine, and "time" falls out as the emergent measure of traversal on a given constraint plane. Process is prior to substance here, which is again Whitehead. There is a subtlety physics forces and the framework absorbs rather than dodges: least action looks *global*, selecting a whole path that extremises an action functional. The globality is an emergent description. The Euler-Lagrange equations it yields are local — at each point the system follows local constraint geometry — and the global path is what emerges when the local equations are solved under initial and boundary conditions, which are themselves constraints in the network. The network does not know the action functional. The action functional is a manifold-level description of what local traversal collectively produces, the way thermodynamic equilibrium is a description of what local molecular collisions produce.

The engine needs a slope to run down, and where the slope comes from answers why there is something rather than nothing. A flat manifold has no preferred direction anywhere — nothing is more satisfied or lower than its neighbours — so nothing changes, nothing stabilises, and no observation is possible within it. Take it as a theorem within the set that a flat manifold is indistinguishable from nothing. The preorder is what distinguishes existence from non-existence: the moment the constraint network carries any directional bias, local traversal is inevitable, and structure follows from it. This also disposes of the standing objection that every system has a preorder, so the claim is vacuous. That mistakes foundational universality for vacuity, the way objecting that every physical system has energy would; the configurations that carry no preorder are not systems, and they are not anything.

4.3 Emergence works by forgetting

Emergence, the last primitive, gets a precise sense so it stops meaning "things got more complicated." Emergence is the appearance of new objects and morphisms by *coarse-graining* — passing to a quotient in which irrelevant distinctions are factored out — such that stable patterns at one level become constraints for the next. Emergence works by forgetting. Renormalization forgets microstates, an interface forgets implementation, a semantic layer forgets tokens; one operation across substrates. The hierarchy is self-generating: constraints yield modes, modes become new constraints, new constraints yield new modes, and at every level the stable patterns are the walls for the next level's standing waves. The hierarchy is also self-bounding, taken as axiom: the chain of emergence does not regress without end downward but terminates at a self-consistent constraint layer whose internal consistency defines what is physically possible — a base that bounds itself with no external wall, supports stable patterns, and generates the manifolds above it. This is the layer Axiom C names as the substrate, and its self-consistency is why the framework needs no ground beneath the ground.

5 The two hardest tests, declared

The framework has to point at where its physics reading is most exposed and least paid off, and name it as such rather than bury it. Two challenges do that work, and both are declared here as interpretive framing that has earned reputational exposure without earning a prediction.

5.1 The proton-electron ratio: an open cheque

The proton-to-electron mass ratio, about 1836, is the first. Any substrate theory that claims constants are emergent invariants of constraint topology owes a derivation of one such constant from first principles, and this is the cleanest target. The honest status is that the framework supplies none. Within the Standard Model neither the QCD scale nor the electron Yukawa coupling is predicted; the ratio is not derivable from the Standard Model alone, and a substrate that genuinely reduced this freedom would be a real result. The framework states the programme — treat the proton mass as known from QCD given its scale, make that scale the first emergent number via renormalization-group flow, attack the Yukawa as a discrete or topological invariant — and concedes that it has zero genuine hits. This is the framework's most expensive open cheque, and it is written as a cheque, not cashed.

5.2 The Higgs as an order parameter

The Higgs is the second, and lighter. In this reading the Higgs is the macroscopic order parameter of a particular constraint phase that sets the stiffness of certain standing-wave modes, and the Higgs particle is a fluctuation of that order parameter; mass is how tightly a standing wave is bound by the global constraint structure, and "giving mass" is the network entering a phase where certain excitations cost more energy to propagate. This is interpretive stance, analogous to phonons in crystals or Cooper pairs in superconductors — effective descriptions, none fundamental in the substrate — and it predicts nothing the Standard Model does not. It is kept because it shows the vocabulary applying cleanly to a hard case, and cut from any claim to be physics.

6 Meeting the rivals

A constructive metaphysics earns its shape partly by where it declines to stand, and the neighbours are close enough that the declining has to be explicit. This section is the adversarial corpus, kept as a real section because reasoning against these programs is where the stance shows its content.

6.1 Einstein and the missing pre-geometric layer

His unification programme was the right instinct aimed one layer too high. He wanted no arbitrary structure — no point particles, no singular sources, no free constants — with all fields and constants emerging from geometry. He tried to build the substrate directly in spacetime geometry, forcing continuous manifolds, singularities, and dimensionful constants as input. The framework's reading is that he was missing the pre-geometric constraint substrate: the layer from which geometry itself emerges. His target was correct; his substrate was one projection too shallow.

6.2 Bohm and the deepest alignment

The deepest alignment, already drawn: implicate order is the constraint network, explicate order is the emergent manifold, unfolding is constraint dynamics projected onto spacetime, particles-as-projections are Markov objects. Where the framework extends him is in trying to characterise the invariants structurally — constraint propagation, Markov objects, the unit of change — rather than through the hologram and the ink-drop, and in carrying the same architecture past physics into computation and engineered systems. Bohm supplied the foundation. This is an attempt to build on it.

6.3 Wolfram and the selection question

The ruliad — the space of all computations from all rules and all initial conditions — is the maximal unconstrained possibility space, and physics is a particular self-consistent constraint manifold within it. The framework diverges on three points. First, there is no single entity called "the constraint network"; there are potentially infinite manifolds, layered and cross-cutting, and the framework identifies invariants across them rather than the properties of one. Second, the observer is not special: Wolfram's programme needs an observer theory in which the observer's computational boundedness determines the physics it sees, while here the observer is a Markov object like any other, and the causal direction is manifold-to-observer — the observer evolved within the manifold and adapted to it, so its properties are already explained by the dynamics that produced it and need no separate theory. Third, and decisively, the selection question is Platonic. Wolfram asks what selects the actual computations from the ruliad of possible ones. That question presupposes the possible exists in a way requiring selection — that the ruliad is a real space our universe is drawn from. This is the same move the framework rejects in Many-Worlds. There is one constraint network; it computes what it computes; the space of all possible computations is our description of what could be consistent, not a real space needing a selection principle. The question "why this universe?" is malformed rather than unanswered, because it posits entities the framework declines to multiply.

6.4 Many-Worlds and uncomputed branches

The principled core is real: take the Schrödinger equation seriously, apply it universally, refuse a collapse postulate, and read apparent collapse as decoherent branching. The objection is not to the mathematics but to the claim that all branches exist. Existence, on the constraint view, requires computation, and this is where computational presentism pays off against a rival. The network computed one outcome at each step; the Schrödinger equation *contains* the other terms as a mathematical description, and Many-Worlds grants substrate-level ontological status to states the network never computed. The disagreement reduces to a choice of ontological primitive — the universal wavefunction containing all branches, or the network's singular computed state — and it is the same error as treating the path integral's sum over paths as evidence that all paths are physically traversed, when the sum is an ignorance computation over a constraint topology the observer cannot see.

6.5 Penrose and the twistor alignment

The Gödel divergence was settled above; the twistor alignment is genuine and worth stating as interpretive stance. Conventional physics separates the stage — spacetime geometry — from the actors — fields and particles — and then struggles to reunify them, which is much of what the quantum-gravity problem is. The framework holds that the separation is an artifact of the projection: at the substrate level topology and constraint are one structure, a single object that is both space and the rules in it. Twistor theory may be the closest existing formalism, because a twistor encodes the null-geodesic topology and the constraints within it as co-defined rather than layered — the incidence relation is a constraint equation, the conformal structure a topological condition. Whether twistor incidence relations can be derived as constraint equations in this framework's precise sense is an open formal question and a concrete bridge worth building.

6.6 Constructor theory: Deutsch and Marletto

The closest formalism to the framework's structural claims: physics stated as which transformations are possible and impossible rather than as equations of motion, with a constructor an object that causes a transformation and retains the ability to cause it again. Possible-and-impossible transformations map to admissible-and-inadmissible morphisms, the constructor to a persisting Markov object, substrate independence to structural invariance. The foundations diverge cleanly — Deutsch's programme is Platonic, treating mathematical structures as independently real, where this framework is Aristotelian — and the divergence does not diminish the correspondence. Constructor theory's treatment of possible and impossible transformations as the fundamental content of physical law may supply mathematical tools the framework's ontology needs.

7 Crossing substrates: the constructor before the encoding

7.1 The recurring encoding-builder pair

Here the same structure begins to precipitate in a second domain, and the reader should watch the physics vocabulary carry without renaming. Every constructive system seems to pair an encoding with a builder: DNA with the ribosome, a specification with a compiler, a prompt with an attention mechanism. The intuitive story runs encoding-first — blueprint, then building. The framework inverts it, and the inversion is the single idea most worth keeping.

7.2 Axiom X: the constructor comes first

Axiom X, the constructor precedes the encoding — the abiogenesis inversion. In information-driven construction the constructor comes first, and the encoding *emerges* as the constructor's solution to reliable replication under constraint. Self-organizing chemistry — the proto-cell — ran before DNA, and DNA arose as that machinery's answer to copying itself faithfully. Informal practice preceded specifications; the builder came first and the spec crystallized as practice's solution to reliable construction. Tradition preceded constitutions. The ordering is `constraint → constructor → encoding → encoding-drives-constructor`, taken as axiomatic within the set. That the ordering is *universal* across every constructive domain — that it must always run this way — is a stronger claim held only as conjecture. Once the encoding exists it takes the driving role, which is why mature systems look encoding-first and hide their own history.

7.3 From flow to representation

The ordering is grounded in Deacon's emergence hierarchy, four steps from flow to representation, and each step happens because Axiom G makes it possible — possibility alone suffices, and need plays no causal role. First **homeodynamic**: self-organizing flow with no stable structure, constraint propagation with no Markov object yet. Then **morphodynamic**: stable structure precipitates from the constraints, Markov objects form as standing waves, crystals, attractor basins. Then **teleodynamic**: a structure begins maintaining its own boundary, a selfreconstructing Markov object, which is the constructor — Deacon's autogen. A Markov object that reconstructs its own boundary is Maturana and Varela's autopoietic unit almost verbatim, which the earlier work re-derived through Deacon without naming them; the condition under which such a boundary holds together is Montévil and Mossio's *closure of constraints* from theoretical biology, and the framework's only addition is to extend closure past life to every substrate that supports it. Finally **representational**: the encoding separates from the mechanism, when the constructor's replication solution becomes a thing you can lift out and copy — DNA, a spec, a constitution. The builder pattern first appears at this fourth step, a late arrival four levels above the substrate, when self-maintaining systems externalize their construction logic.

7.4 Inhabitable rather than merely possible

This is where Kauffman and, more sharply, Cronin and Walker's assembly theory belong. The question Axiom X poses — what, among the possible, is *inhabitable*; what a constraint landscape will actually build and select for — is assembly theory's selection question, and the assembly index is a live candidate measure for the constructor-first ordering. It is also the division of labour that separates this programme from universal computation: universal computation asks what is possible, and the constraint-emergence programme asks what is inhabitable. Stated as a pattern the reader can carry across domains, `encoded representation → constructor → constructed structure` reads:

Domain	Encoded representation	Constructor	Constructed structure
Biology	DNA	ribosome / cellular machinery	protein / organism
SDLC	specification (requirements, schemas, contracts, tests)	builder (LLM agent, compiler, developer)	artifact (code, config, deployment)
LLM	prompt / loaded constraint spec	attention mechanism	evaluation / output
Neuroscience	brainstem priority structure (affect)	cortical computation	behaviour

Each row is the same four-step climb reaching its representational stage in a different substrate. The table summarizes the tour; it does not replace it, and the authority stays with the prose.

8 Where meaning enters: the evaluator

8.1 Evaluation is the threshold

A constructor builds, but building toward what, and where the target comes from is where meaning enters the ontology. The threshold for meaning is *evaluation*, and life and intelligence both sit far above it. Something has to hold a model of a target, notice the gap between where things are and where they should be, and push.

8.2 Axiom E: the evaluator emits direction

Axiom E, the evaluator emits a constraint signal. Meaning enters exactly where a Markov object *evaluates*: models a target, computes a delta between current and target state, and emits that delta as a constraint signal directing a downstream traversal. The evaluator is a role — any Markov object that does this is one — and one architecture recurs at every scale: `evaluator → constraint signal → constrained traversal → feedback → encoding update`. Solms's work on the brainstem gives the concrete biological instance. The reticular activating system does not perform detailed cognition; the cortex does. What the brainstem does is evaluate danger, reward, and novelty and emit the result as affect that constrains what the cortex does next. Emotions are prompts. They set the preferred direction within which the cortex then traverses its own manifold, and the cortex operates under the constraint without needing to understand why the brainstem evaluated as it did. This is structurally identical to a human prompting a language model: the human evaluates, produces a constraint signal, and the model traverses its semantic manifold under it — the prompt is the context parameter of the preferred-direction function, not the computation. Stated in the operator's own terms, this is the supervisor loop said ontologically: the ticket is the emitted constraint signal, and the model-plus-telemetry-to-gap-to-intent cycle is one turn of Axiom E, which is why the loop runs ticket-first.

8.3 Meaning is the invariant operation

From this a theorem within the set falls out: meaning is a structural invariant. The meaning-operation — a pattern matched against a model, match or no match — stays fixed at every scale. A thermostat matching temperature against a threshold, a moth matching light-angle against a hard-coded rule, a cat matching a trajectory against a predictive model, a human matching a narrative against a simulated counterfactual: the same invariant run into increasingly rich constraint spaces. All the richness lives in the space; the operation is constant. A more complex evaluator grows the manifold that meaning maps into. Meaning itself does not scale, deepen, or grow richer. The floor for meaning is the simplest system that represents and matches. A crystal settling into an energy minimum sits below that floor — it obeys, with no model and nothing represented, only constraint dynamics.

9 The observer is a Markov object

Now the same structure precipitates a third time, in the domain where the classical puzzles of physics live, and it dissolves one of them. The word "observer" conflates two unrelated things, and the conflation generates the measurement problem.

9.1 Two kinds of observer

There is a **Type-1** sense: measurement as a physical interaction between constraint patterns on a manifold. A photon hits a detector; a particle scatters off another. This is constraint propagation, patterns interacting and the constraint structure resolving, with no special agent and no observation in any ontologically meaningful sense. This is physics. There is a **Type-2** sense: the evolved observer, a brain, a computational predictive engine refined over hundreds of millions of years, a Markov object of extraordinary complexity that models its environment by running internal constraint satisfaction. It is special in complexity. It is not special in kind: a stable pattern on an emergent manifold, distinguished only by the depth of its internal constraint structure.

9.2 Why the measurement problem dissolves

The measurement problem asks why observation produces a definite outcome. The question dissolves once the conflation is separated. Type-1 interactions require no observer; the constraint structure resolves regardless of whether any evolved agent is present. Type-2 observers see one outcome because they are computational structures on an emergent manifold that represents the computed state. The observer accesses the information available at its layer. It does not cause the outcome. The causal direction is manifold-to-observer. The observer is a product of the constraint dynamics it inhabits, adapted to the manifold's structure through evolution, itself a constraint-satisfaction process over deep time; its perceptual categories, predictive models, and sense of locality are consequences of the manifold it evolved within. Any theory that begins with the observer's properties and derives physics from them has the causality backwards. The observer needs no special theory, only an evolutionary account, which is constraint dynamics applied to self-replicating Markov objects.

9.3 Meaning and caring emerge separately

Meaning and caring are distinct emergences on distinct manifolds, and keeping them distinct matters because the framework grants meaning to non-living evaluators. Meaning arises wherever a Markov object evaluates. Caring arises at Deacon's teleodynamic transition, where constraint closure becomes self-maintaining and the Markov object begins to act on its own behalf — persisting, repairing, reproducing, acting to maintain its own boundary conditions. A language model's outputs mean something; the model does not care. An organism's signals mean something; the organism cares. The meaning is the same kind of thing, evaluation producing reference; caring is a different kind of thing, arising on the life/non-life manifold rather than the meaning/no-meaning one.

9.4 The gradient is model depth

The invariance of meaning becomes visible along a gradient of model depth, which is Solms's contribution read across biology. A moth circling a light runs a hard-coded rule with a minimal internal model; its computation is distributed almost directly back into the environment, and the gap between stimulus and response is small. A cat tracking prey maintains a richer model — it predicts trajectories, waits, adjusts — and the gap widens, evaluation happening within the model before action returns to the world. A human operates almost entirely within a simulation of reality, a vast internal model of counterfactuals and narratives; emotions exist as the mechanism that prompts this simulation, the brainstem's evaluation breaking through to redirect the cortex as a hard interrupt, because the simulation is so dominant that redirection needs one. Beneath it all the hard-coded reflexes persist, the insect-level architecture still running, still bypassing the simulation when speed matters more than accuracy. The meaning-operation is identical at every point on the gradient. What differs is the constraint space the evaluation matches against.

9.5 Structural words without anthropomorphism

A note on method, since these are the words most easily misread: throughout, terms saturated with human qualia — meaning, intent, emotion, caring, observer — are used deliberately and then stripped to their structural operation. The qualia-rich word draws on the reader's experience; the functional definition reveals the invariant. The felt quality of meaning is real, and it is what pattern-matching looks like from inside a human simulation, one manifold's projection of the invariant rather than the invariant itself.

10 The working law: expansion needs collapse

The structure has now precipitated in physics, in biology-and-construction, and in the observer. The consolidation is a single law that turns the picture into a design discipline, and it is where the ontology stops being contemplative.

10.1 Expansion needs a contracting boundary

Watch what Axioms T and M say together. By Axiom T a probabilistic generator traverses a constraint manifold and *expands* — it opens into the whole set of admissible continuations. By Axiom M a stable output requires a boundary that screens, which means the expansion has to be over-determined down to one configuration: a collapse. Something has to supply that over-determination. Follow it through and it is a theorem.

10.2 Theorem H: the F_P to F_C regime law

Theorem H — expansion→collapse is the F_P→F_C regime law. A probabilistic generator (F_P) expands into admissible continuations; a stable output requires a boundary that collapses that expansion to one configuration; a contracting verification contract (F_C) supplies the collapse. Therefore probabilistic compute without real contraction is hallucination — expansion with no collapse leaves the traversal in a degenerate region where many continuations carry equivalent probability and none is grounded. F_P expands, F_C contracts, and stable emergence needs both. This is the same shape as superposition→decoherence and as mutation→selection: a proposer under a disposer. One reprice from the earlier working vocabulary is ratified in the naming. The contracting regime first travelled as F_D, deterministic contraction, and determinism turns out to be one implementation property of a contractor rather than the invariant role — a stochastic test can supply real evidence, and a deterministic program can be wrong with perfect repeatability. So the theorem's statement carries the repair: verification is a contract among roles, deterministic checking is one role inside that contract, and it is never truth itself. The status is the point of the whole discipline. Theorem H is *derived within the set*, earned from the axioms, which is a different and weaker kind of claim than a stake against the world, and the framework keeps the two from ever blurring. It is here that the collapse of the physics section pays off — collapse is interpretive stance when it reads decoherence and a theorem within the set as this compute regime, and the status law lets the one word carry both without wearing both at once.

10.3 Four roles keep assurance honest

Look closely at that contract and four roles precipitate, each already latent in the axioms. A **proposer** produces candidates, which is F_P wearing its role name. An **evaluator** scores candidates against a criterion, which is Axiom E's Markov object emitting its delta. A **verifier** checks a declared property and returns evidence under a stated contract — a test suite, a type checker, a measurement, a proof. An **admitter** decides which candidate becomes accepted state, publication, or action, and authority lives there. One component can implement several roles; the contracts never merge. A compiler verifies the properties it was built to see and admits nothing. A reward model evaluates and verifies nothing. A human or a policy gate admits, and its authority is granted by the surrounding system rather than generated by any check it runs. Four assurance properties keep the roles apart, and none implies the others: determinism asks whether the same admitted inputs reproduce the same result; correctness asks whether the declared property actually holds; grounding asks whether the check is tied to execution, measurement, or formal rule with provenance; selection authority asks who may make the result operative. A check can be deterministic and wrong, correct and ungrounded, grounded and unauthorized. Take the separation as a further theorem within the set, falling out of Axiom E the moment the evaluator is placed inside a world rather than inside a closed score.

10.4 One hybrid manifold

The two regimes compose into one hybrid manifold. **F_P**, probabilistic expansion, is stochastic generation of candidates from the encoding — model generation, design exploration, retrieval-augmented proposal — traversal that widens the admissible set. **F_C**, contraction, is the verification contract that over-determines — tests, type checks, schema and lineage contracts, policy gates, admission rulings — traversal that collapses the admissible set to what survives, and **F_D** names the deterministic checking role inside it. So the framework yields a design law for anyone building constraint-based world models: build the contracting boundary explicitly, roles and authorities declared, and the system inhabits stable Markov objects; omit it and the system hallucinates. This is the mapping onto the operator's own F_P/F_D evaluator regimes, and it is genuine rather than flattery — his line "F_P without F_D is hallucination" is the same claim as "sparse constraints produce degeneracy," his approved typed artifacts behind interfaces are Markov objects, and the reprice sharpens his law rather than retiring it, since the deterministic gate stays mandatory and stops claiming to be the whole of truth's contract.

10.5 Truth costs a grounded collapse

The regime law gives a working definition of truth, and the definition is what the harness later grades transfer against. Take it as a theorem within the set that truth is a stable Markov object in a constraint manifold — the configuration that survives a grounded collapse and resists perturbation. Stability alone never confers it, and the role split says why: a false belief can be institutionally stable, a model can hold a deep attractor for an obsolete fact, a program can pass every test that shares its blind spot. In each of those cases something persists, and it persists because a role in the contract was missing — the check ungrounded, the admission unearned. In physics a truth is a standing wave that persists under perturbation. In a language model it is an attractor basin that stays coherent under resampling and verification. In software it is an artifact that passes all constraints and holds under selection pressure. What F_C supplies, in each substrate, is the collapse that leaves a truth standing where an ungrounded expansion would have left a degenerate region. Truth is not a property a proposition has apart from the manifold; it is what the manifold stabilises when the collapse is real, and the four roles are what “real” costs.

11 The hallucination spiral, in five passes

The hallucination point is load-bearing enough to earn a full spiral, each pass adding a distinct layer, and the earlier work’s error was restating it flat some five times rather than deepening it. It is walked once, properly, here.

11.1 Pass one: the phenomenon

In manifold regions where Markov objects did not form during training, the probability landscape becomes degenerate: multiple continuation paths carry equivalent probability. The model keeps traversing with full confidence, because from its perspective every path is equally valid. The label “hallucination” belongs to the observer. The model cannot distinguish constrained from unconstrained traversal, and the output is fluent and ungrounded not because the trajectory is unstable but because the constraint surface cannot discriminate between paths.

11.2 Pass two: the physical identity

This is structurally the same as quantum indeterminacy. In a degenerate eigenspace — multiple eigenstates, one eigenvalue — the system has no structural reason to prefer one state, and the outcome is fixed by measurement rather than by the dynamics. At the topological level the equivalent paths are homotopically indistinguishable: the constraint surface lacks the information to discriminate them. Hallucination, quantum indeterminacy, and topological path-equivalence are one phenomenon across substrates — the manifold is degenerate and the resolution is always external to the system.

11.3 Pass three: the engine mechanics

By Axiom T the model does not know its destination; it evaluates the local preorder and moves in the preferred direction. Stable reasoning is laminar flow into a deep attractor basin, a Markov object. Hallucination is what traversal does when the landscape is too sparse to carry a preferred direction — the manifold flattens, the preorder gives out, and any continuation is as downhill as any other. This is why grounding helps and why prompting works: external constraints add boundary conditions that restore the preorder and stabilise the attractor, and prompting reshapes the local constraint geometry so that different Markov objects can form.

11.4 Pass four: the owned engineering prediction

In software the same condition is sparse specification. Artifacts fail when built from too few tests, unclear requirements, or novel architecture with no prior construction history to inform the encoding — the specification is insufficiently constraining, multiple constructed structures look equally valid to the builder, and the resulting artifacts pass local checks but fail under real selection pressure in deployment. The prediction is specific and owned: failure rates should correlate with specification density, not with code complexity or team size as such. This is the SDLC image of the same degeneracy, and it is testable in the operator’s own system.

11.5 Pass five: the formal statement

One typed phenomenon has three images — a degenerate eigenspace in physics, a flat logit region in a language model, an under-specified requirement in software — and it has a formal name in the cost structure of the next section: degeneracy is the condition where constraint density is too low to collapse the expansion. Dense constraints collapse cleanly. Sparse constraints leave the manifold degenerate.

12 The cost of traversal

12.1 A Hamiltonian for traversal

Traversal has a cost, and the cost has a shape — a further theorem within the set. Each step composes a traversal Hamiltonian, `H(step) = T(iteration_cost) + V(constraint_delta)`: a kinetic term T, the cost of one compute step — metabolic energy, compute cycles, human hours — and a potential term V, how far the current state sits from satisfying every constraint. Constraint density is the metric on the manifold: it sets how far apart states are and therefore how expensive a transition is. Dense constraints mean low V, cheap and clean convergence; sparse constraints mean high V, degenerate and expensive, which is the very condition that produces hallucination. Three things that looked separate are one: constraint density sets local update rates, sparse constraints breed degeneracy, and convergence cost scales inversely with density. The ground state is the minimal-effort configuration satisfying all constraints; over-engineered solutions are excited states carrying more energy than necessary; optimal trajectories minimize ΣH. Least action turns out to be a universal property of constraint-manifold traversal, of which its appearance in physics is one image, and its appearance in a build loop, a metabolism, and an inference pass are others.

12.2 Intent lineage is the composite trajectory

There is a structural corollary worth stating, because the harness depends on it. A complex output is a composite trajectory through the manifold: each intermediate is a Markov object that becomes a constraint for the next, which is Axiom E turning over and over. The composite carries its full **intent lineage** — the traceable chain from the first constraint signal to the final stable structure — and any edge in that chain is zoomable, expandable into a finer sub-graph or collapsible into a single edge, because the same four elements (manifold, local traversal, evaluator, encoded representation) apply at every scale.

13 The single vocabulary, as a map

13.1 Six terms across four domains

Everything above uses six typed terms, each with one meaning, and the payoff of the discipline is that the same six carry across every domain the framework touches. Each cell is an asserted *image* of one typed term under a domain-specific map; whether the maps are faithful is the question the next section takes up, and the table leaves it open.

Typed term	Physics	Biology	LLM	Software (SDLC)
constraint	admissibility of a transition	admissible reaction / selection condition	prompt + learned distribution over next tokens	requirement / test / type / contract
Markov object	standing wave, bound state, particle	cell, organism (autopoietic unit)	attractor basin / proto-symbol	approved versioned artifact behind an interface
boundary	potential well / interaction surface	membrane / organizational closure	basin separatrix in activation space	interface / API contract
collapse	decoherence, measurement	commitment to a phenotype/pathway	sampling a token from the logit distribution	UAT approval + deploy
traversal	least action (local Euler–Lagrange)	metabolic/developmental dynamics	forward integration of the direction field	builder reducing the gap to done
emergence	renormalization / effective theory	morphogenesis, teleodynamic transition	layer hierarchy as abstraction levels	composition of artifacts into higher assets

13.2 Images are not proven identities

The claim the table makes is exactly this and no more: each row lists images of one typed term under domain-specific maps. Calling something an "instantiation" here means asserted image under an as-yet-unproven map — a license for the loaded model to carry structure across domains during exploration. Proven identity is a different claim, and the table never makes it. This is the single most important honesty the cut enforces, because the earlier work's credibility flaw was exactly here: its prose insisted "this is not analogy, it is the same operation" while its formal section conceded the map was unproven. Under this constitution, "treat these as the same operation" is legitimate as an instruction to the loaded model to transfer structure and see what generates, provided it wears its instruction status openly. Dressed as a discovered theorem, the same sentence becomes the old flaw again.

13.3 Every crossing declares fidelity and loss

The license carries a declaration law, and the law is constitutional. Every crossing made with this vocabulary owes six declarations at the place the crossing is made: the source domain, the target domain, the relation claimed preserved, the fidelity with which it is preserved, the loss the crossing accepts, and the condition under which the map fails. A membrane and an API contract both mediate what crosses a boundary, and they differ in metabolism, materiality, authorship, and authority; an atom and an approved artifact both persist, and only one of them has physical dynamics. Declaring the loss is what keeps the shared word honest; declaring the failure condition is what keeps the map a claim rather than a mood. And where a declared map breaks, the breakdown is a result in its own right, because it locates the border between the invariant the framework proposes and the mechanism the substrate supplies.

The framework therefore owes a place where the maps stop being asserted and start being able to fail.

14 The formal spine and the two exposure points

14.1 Why the category comes last

Now, and only now, the category is named. A constitution that only ever reasons from its axioms is unfalsifiable in the empty way that never touches the world; a metaphysical ultimate holds a license for its unfalsifiability, but a whole framework floating free of data holds none. So the framework stakes two claims against outside data, and it builds them where the substrate is one we can inspect — the language model, whose internals mechanistic interpretability can open.

14.2 Category C and its two functors

The formal spine is a category, and here we are in the register of conjecture, a structure proposed and awaiting proof. Let **C** be a category whose objects are Markov objects `(S, B, D, σ)` — state space, boundary operator, dynamics operator, stability condition — each carrying the Markov property, and whose morphisms are the boundary-, screening-, and stability-preserving maps that compose and admit identities. Two morphism classes matter: *emergence*, taking components to a composite whose boundary screens their internals — atoms to a molecule, token features to a concept attractor — and *collapse*, taking a superposition to one object under over-determination — wavefunction collapse, context resolving a polysemantic feature to a single meaning. Two functors are asserted out of C — `F_phys` into physical bound states and `F_llm` into activation space. The honest status split is that `F_phys` largely draws on established Markov-blanket results and is close to done, while `F_llm`'s faithfulness is the open problem. Axiomatizing C and proving that composition preserves conditional independence is the immediate formal task, and it is legitimate mathematics inside emergence, distinct from the category error of demanding a substrate proof.

14.3 Why training may produce Markov boundaries

The reason to believe the LLM functor could be faithful is the structure of training, and this argument is derived within the set and feeds the exposure point that follows. Gradient descent on structured data makes semantically similar inputs activate similar patterns — attractor basins; transitions between basins pass through unstable regions — the boundary operator; a cleanly represented monosemantic feature keeps consistent internal activations regardless of distant context — conditional independence given the boundary; polysemanticity is *incomplete* screening, boundaries overlapping when Markov objects share neurons so external context leaks; and scaling adds dimensions that support cleaner separation, so screening should improve with scale. The existence of these regions is settled by training. Only their precision is empirical, and that precision is the first place the framework can lose.

14.4 Exposure point 1: the Conditional Independence Conjecture

This is a stake against outside data. For a trained LLM feature identified by a sparse autoencoder, with internal dimensions I, boundary dimensions B, and all others E, we stake that `P(A_I | A_B, A_E) ≈ P(A_I | A_B)` to within a tolerance ε that *decreases with model scale.* The tests are concrete: ablate features outside the boundary and watch internal activations move by less than ε; measure ε across model sizes and watch it fall; check that monosemantic features screen better than polysemantic ones; vary distant context and confirm boundary-screened features stay low-variance. It falsifies cleanly — if ε does not fall with scale, or monosemantic and polysemantic features screen equally, the screening property is not a consequence of training and the functor claim collapses to loose analogy. The honest scoping matters: this conjecture's content sits largely downstream of existing mechanistic-interpretability work, it is a stake built *on* that literature, and it should wear that provenance openly.

14.5 J-space is a hypothesis generator, not a proof

One adjacent piece of evidence gets its scope fixed here, because it is the easiest to over-read in the conjecture's favour. Jacobian-lens work on trained models finds a small, causally privileged, verbalizable workspace — J-space — that supports report, directed modulation, and flexible reuse while much automatic processing bypasses it. The framework takes that as convergent evidence for a privileged workspace, and takes it as nothing more. A broadcast workspace makes selected information widely available; a Markov boundary screens an interior from an exterior conditional on the boundary state; the topologies differ, and evidence for the first never carries the second. J-space is admitted as a hypothesis generator — a place to look for boundaries worth testing — and the conjecture still wins or loses on the screening tests themselves.

14.6 Alignment as a commuting diagram

The functor also lets us say precisely what "the model models the world" means, which gives alignment the form of a diagram, though this too stays conjecture until η is exhibited. A world model is a functor from a schema category — what entities exist and how they relate — into the LLM category — which basins represent which entities, which transitions represent which relations. The correspondence between model and world is a natural transformation η between `F_phys` and `F_llm` out of that shared schema, mapping each physical Markov object to its LLM representation while preserving morphism structure. **Alignment is the question of whether η commutes** — a misaligned model is one where η fails at identifiable morphisms, where the model's internal structure diverges from the world's at nameable points. Explicit world-model architecture — tool boundaries, planning states, chain-of-thought — makes the otherwise-implicit Markov objects observable, which is what makes η testable. In the same spirit the functor would generate paired predictions for derivation: energy minimization giving stable bound states beside loss minimization giving stable basins; phase transitions beside capability jumps at scale thresholds; the renormalization group beside the layer hierarchy as an abstraction flow — each a candidate still to be earned, none a resemblance to be banked.

14.7 Exposure point 2: the owned world-model claim

The second stake against outside data. Constraint-based world models built with *explicit* Markov-object boundaries — typed assets with declared interfaces — show measurably lower hallucination and better out-of-distribution screening than matched baselines without explicit boundaries. This is the framework's most directly owned stake, because it is testable directly in the operator's own system with frontier-model access optional, and the status difference the discipline forces is exact: Theorem H says "F_P without F_D is hallucination" and is derived within the set, while this claim is staked against outside baselines and can be beaten by them.

15 The epistemic apparatus

The framework's honesty is itself structured, and the structure is part of what the harness grades. Three instruments carry it.

15.1 The status ledger

The first is a status ledger. Every substantive claim in this document carries one of the five statuses, and the load-bearing ones can be read off directly: the constraint-network substrate, the generative principle, the Markov-object units, traversal-as-engine, the constructor-before-encoding ordering, and the evaluator role are **axioms** — the framework assumes them, and rejecting one exits the framework. Computational presentism, the Gödel/Turing/Cantor unification, meaning-as-invariant, Theorem H with the four-role separation it carries, the traversal Hamiltonian, and least-action-as-universal are **theorems within the set** — earned by derivation, challengeable only on the derivation. The universality of the constructor ordering, the faithfulness of `F_llm`, the category C, and the natural transformation η are **conjectures** — open inside the set. The map/territory reading, the physics interpretations of collapse and motion and gravity and the Higgs, stratified time, and the twistor alignment are **interpretive stances** — compatible with the physics, carrying no independent test, to be used as lenses and never asserted as discovered fact. And the screening precision of Markov objects and the owned world-model advantage are **empirical-exposure-points** — the two places the framework stakes against outside data and can actually lose.

15.2 What would falsify the framework

The second is falsifiability with its complement declared. A framework that cannot be falsified is not useful even as philosophy, and the following would damage or destroy it: a proof that no structure-preserving map exists between the constraint structures of any two domains, which would collapse the structural-realist thesis to loose analogy; a derived physical constant that contradicts observation; a demonstration that LLM attractor basins do not exhibit conditional independence, which fails the most testable claim; a system where stable emergent structure arises with no boundary, symmetry breaking, or constraint geometry, which would break the absential-causation engine; or a rival ontology that makes all the same structural claims and also generates unique predictions this one cannot. What would *not* falsify it is failure to derive new physics — the framework positions itself as pre-physics, and remaining philosophy forever is stagnation rather than falsification, a different and lesser failure mode, but the one this cut's open harness is meant to prevent.

15.3 Evaluation modes and their costs

The third is a set of legitimate evaluation modes, each with a cost if it fails, because the framework changes what counts as evidence. Computational genericity — toy constraint-network programs robustly producing persistent localised modes and effective continuum limits — is partially achieved and its failure would weaken the generic-engine claim moderately. Cross-domain structural transfer — a rigorous structure-preserving map between LLM activation geometry and toy physical constraint propagation — is open and its failure would be very costly, collapsing the correspondence claim. Interpretability signatures — clean conditional independence at boundary regions in frontier models — are ongoing and patchy and are the strongest near-term anchor. Failure to coarse-grain toy models to known physics would be moderately costly. Deriving even one exact dimensionless ratio remains at zero genuine hits and its permanent failure would be the most expensive, because it is the strongest unique prediction. If several of these keep returning negative over three to five years the framework should lose credibility even among the sympathetic; if a critical mass return positive, treating mathematics as ontologically prior starts to look prohibitive, which is the point. The research directions compress to the same order the functor programme implies: define C and prove composition preserves conditional independence, and exhibit `F_phys` formally, both doable now; test the Conditional Independence Conjecture with sparse autoencoders and causal tracing across scales, doable with current tools; formalize world models as functors and test whether η commutes once explicit architectures exist; and at the far end derive a known scaling law from C and design architectures with enforced Markov boundaries, where validation turns into engineering payoff.

16 How a cut proves itself

16.1 A constitution closes by experiment

Which brings us to the destination, and to the reason this document is a constitution. Under its governing intent the framework's success is something a run either shows or fails to show; argument has no purchase on it. A constitution proves itself by what emerges within it.

16.2 The matched comparison

So the harness is the paper's real closing argument, and running it is what closes a cut. The experiment is a matched comparison: **with-ontology**, this artifact loaded as context, against a **baseline** that is the same model on the same tasks without it, or against a scrambled-vocabulary control holding token count constant, so the measurement isolates the structure from the sheer length. The judgment runs on three axes, each with a rubric. *Within-set coherence*: do the outputs stay consistent with the axiom set and use the typed vocabulary without drift — every substantive claim carrying a valid status marker, no term used outside its defined meaning, no derivation contradicting an axiom? *Cross-domain transfer*: does the model carry structure across physics, biology, LLM, and SDLC in ways that survive the isomorphism-claim test — a named map that holds under adversarial challenge, where an equivocation dressed as a discovery fails? *Novelty*: does the loaded artifact produce candidate ideas the baseline lacks, that a separate reviewer judges sound?

16.3 Task battery, lineage, and regression

The protocol exercises each cut on a fixed task battery — exploratory, adversarial, applied, derivational, and cross-domain prompts — graded by a judge model against the three rubrics and, for the exposure points, against outside data. Because a complex output is itself a composite trajectory through the manifold, each intermediate a Markov object that becomes a constraint for the next, the harness records the full intent lineage of every run: the traceable chain from the first constraint signal to the final output. Lineage is what makes a run auditable and a regression locatable, and it is the same ledger discipline the operator applies to world-model cuts, brought inside the ontology's own evaluation. Method discipline follows: closure is by experiment alone, and a cut that improves the prose without moving the three axes has changed nothing the harness can see. Scores are recorded against the cut number so regressions surface across the ledger. This cut, v2.0, predicts higher within-set coherence than the untyped v1 and richer cross-domain transfer than the intermediate compressed working cut, since it restores the full traversal without reopening the term-overload defect — and that prediction is itself a measurement this harness is built to take, which is the first thing v2.1 should do.

16.4 The negative record travels with the cut

The harness can be trusted because the framework has already let itself lose, and the record of the losing travels with the cut rather than staying in a lab notebook. The Markov-object experiments bearing on the first exposure point have run, in open models from GPT-2 small through Pythia-160M to Llama-3 8B, and the record is mixed in the way that makes a ledger worth keeping. The sparse-autoencoder partitions leaked: overwriting features outside a target's active set still transferred roughly a quarter of its identity, so the dictionary sensed and fragmented semantic structure without isolating a blanket. Native residual directions were layer-sensitive and causally manipulable — a rank-one direction at an early layer cleared its preregistered criterion, composition was lawful with a compound direction aligning to the sum of its constituents far beyond chance, and modest transfer replicated across a second architecture — while every screening gate failed: with the chosen identity subspace removed, a probe still read the identity perfectly, and five static formulations plus one dynamical formulation failed at both tested scales. The control gives the record its edge: an explicit-object toy transformer passed the same instrument, so the instrument can see a boundary where the architecture actually contains one. The disposition is exact. The negatives demote the tested instruments; they leave the conjecture standing, because a failed instrument-object pairing is a demotion and never a proof of absence; and no later narrative converts a failed gate into a pass without a new preregistered hypothesis. A harness whose ledger already carries recorded losses is a harness whose future passes mean something, which is why this record is constitutional content rather than optional commentary.

16.5 What closure means

The framework asks to be loaded, computed over, and measured; belief is beside the point. It stakes, on its own harness, that a model reasoning inside these six typed terms and seven axioms will explore more coherently, transfer more faithfully, and invent more usefully than a model without them. If that is true, the constitution has proven itself in the only way a constitution can, by what emerges within it. If it is false, the cut does not close, and the next one is owed.

---

17 Cut Digest — v2.0

17.1 What this cut is

A ratified merge. The narrated unified cut is the constitutional base of this file — its register ruled correct by the owner — and a parallel flattened rewrite of the same material served as a content source, mined rather than retained. Four deltas were harvested from it and woven in as narrated prose. First, the four-role verification contract — proposer, evaluator, verifier, admitter — with the F_D→F_C reprice, woven into the regime-law section; it repairs the overclaim that deterministic verification is truth, and Theorem H's statement now carries the repair: verification is a contract among roles, and deterministic checking is one role inside it, never truth itself. Second, the negative Markov-object experiment record — leaked SAE partitions, failed static and dynamical screening gates, the passing explicit-object toy control — woven into the closing harness section as the record that makes the harness honest. Third, J-space scoped as workspace convergence and never blanket proof, placed beside the first exposure point. Fourth, the mandatory fidelity, loss, and failure-condition declaration on every cross-domain crossing, with the physics interpretations statused per projection, reinforcing the vocabulary law and the traversal. The register ruling travels with the merge: the narrated voice, the inline five-status law, the question-bridge transitions, and the precipitation order are constitutional for this paper; numbered spec headings, status tables, and code-block contracts are not admitted as its register. Descriptive third-level headings expose the argument's existing turns for print navigation without changing that register, order, or authority.

17.2 Where the base came from

A unification. v1 was the long-form cumulative traversal — stance, then physics as vocabulary factory, then rivals, then hardest tests, then substrate crossings, then formal functor — but it let its central term mean six things, asserted identity its own formal section conceded was unproven, and never operationalized its success criterion. An intermediate compressed working cut fixed the vocabulary, the status discipline, and the harness, but flattened the physics and lost the traversal that lets a reader watch the structure precipitate. This v2.0 restores v1's arc inside that typed, status-marked, harness-closed frame: the single full paper.

17.3 Restored from v1.3 and reworked into the v2 voice

The derivation of computational presentism from the discrete-step axiom, carrying the "what exists is what has been computed" claim as a theorem rather than a slogan. The Gödel/Turing/Cantor unification as one recursion result, with the Penrose non-computability rebuttal and the substrate-versus-between-domains category-error defense. The full physics vocabulary factory: the hydrogen-atom Markov object, the potential-well boundary, the three-way decomposition of collapse, wire-drift traversal, the nerf-ball rendering of absential causation, and stratified time — each teaching a typed term concretely before it travels. The observer theory: Type-1/Type-2 conflation dissolving the measurement problem, meaning-versus-caring, and the Solms meaning-gradient with emotions as hard interrupts. The adversarial corpus as a real section: Einstein's missing pre-geometric layer, the Bohm mapping, Wolfram's three divergences including the Platonic selection question, the Many-Worlds computation critique, the Penrose twistor alignment, and constructor theory. The hallucination point as an explicit five-pass spiral. The spec-density failure-rate prediction and the why-something-rather-than-nothing /

least-action absorption. The epistemic apparatus: status ledger, falsifiability with its complement, and evaluation-modes-with-costs. Parts III and IV restored as short, declared hardest-tests.

17.4 Kept from the compressed working cut

The computation-over-constraints premise early; inline five-status epistemic marking under a compressed Reading Contract; the honest one-strand-of-physics drift-opening; direct assertion with denial placed as its own sentence where load-bearing, no single-sentence “not X but Y” pivots; typed vocabulary with one meaning per term; the Simondon/Whitehead/Deleuze and Montévil-Mossio/Varela/Kauffman/Cronin-Walker lineage woven inline; both empirical-exposure-points; the evaluation harness as closing destination; the versioned-cut law. These are non-negotiable.

17.5 What remains excluded, and why

The Edinburgh / Foucault / game-theory aside is dropped: it earned nothing under any framing and the review flagged it as a political digression. The variable-c precision-test section, the Smolin divergence table, the cosmological-constant and gravity-entanglement research items, and the anomalies-as-signals list are cut to their single load-bearing sentences or dropped, because they multiply reputational exposure without earning predictions and belong to the physics-mechanism elaboration this framework declines. The bidirectional “prediction” tables that paired phenomena by resemblance (hallucination beside vacuum fluctuations) are dropped as equivocation dressed as derivation; only the functor-generated paired predictions survive, and only as candidates to be earned. The 24-principle consolidation list and the Feynman-diagram extrusion reading are dropped as restatement. The standalone Emergent Reasoning appendix is folded by reference into the LLM functor argument rather than repeated.

17.6 Declared closure gap

The evaluation harness is specified but not executed. This cut is not closed until it runs. v2.1’s first task is to measure the three axes, with the standing prediction that the restored traversal raises cross-domain transfer over the intermediate compressed cut while holding its within-set coherence gain over v1.