

**Emergent Reasoning in Large Language Models:

A Topological and Constraint-Based Formalization (Version 2)**

Abstract

Large Language Models (LLMs) increasingly demonstrate reasoning behaviors that appear to transcend their probabilistic neural foundations. This paper proposes a unified formal framework that explains these behaviors through *topology-constrained traversal* within a learned semantic manifold. We reinterpret the Transformer architecture as a dynamical system whose core computation is a *preferred direction function*, implicitly implemented by attention. Context induces constraint sets; embeddings give rise to conceptual topology; attention performs soft, graded unification analogous to symbolic unification; and trajectories over the manifold follow structured flows that manifest as reasoning.

We show how clusters of embeddings form proto-symbolic attractor regions with Markov-blanket-like boundaries, enabling symbolic-like behavior to emerge from continuous computations. Hallucinations arise when trajectories exit well-constrained regions and enter weakly structured attractors. We compare this with Prolog’s symbolic unification and show that both paradigms instantiate a shared abstract unification principle—one discrete and explicit, the other continuous and implicit.

Finally, we outline architectural mechanisms for transforming emergent reasoning into *explicit, controllable, modular reasoning*, bridging connectionist and symbolic AI. This framework provides a principled understanding of how probabilistic systems can yield structured inference, offering a foundation for next-generation neuro-symbolic architectures.

1 Introduction

The emergence of reasoning-like behavior in Large Language Models challenges traditional dichotomies between symbolic and connectionist AI. These models perform tasks such as:

- multi-step reasoning,
- analogical inference,
- constraint satisfaction,
- factual and relational reasoning, and
- structured problem solving,

despite containing no explicit rules, logic engines, or symbolic manipulators.

This paper argues that the explanatory gap can be closed by reframing LLM computation as *constrained topological traversal* in a learned semantic space.

1.1 A Constraint-Based Perspective on Reasoning

Classical constraint satisfaction problems (CSPs) reduce solution space via symbolic constraints: [$P' = \{ x \in N : C_1(x) \wedge C_2(x) \wedge \dots \wedge C_n(x) \}$.]

Similarly, in an LLM:

- the input prompt imposes contextual constraints,
- attention weights dynamically enforce relevance constraints, and
- learned embeddings narrow possible semantic continuations.

The Transformer thus performs *constraint-directed computation*.

1.2 LLMs as Dynamical Systems over a Semantic Manifold

We introduce a formal model where:

- (M^d) is a **semantic manifold**,
- trajectories $(x_0, x_1, \dots, x_T \in M)$ represent **semantic evolution**, and
- the Transformer implements $[x_{t+1} = F(x_t, c_t),]$ where (c_t) is context.

The core computational primitive is: [$D : M \rightarrow T(M)$,] a **preferred direction function**, implicitly realized by attention.

1.3 Soft Unification and Emergent Symbolic Structure

Attention performs weighted matching via similarity: [$(QK^H)V$.] This constitutes *soft unification*: the continuous analogue of Prolog's logical unification.

Embedding clusters form attractor-like regions with Markov-blanket-like conditional independence properties. These **proto-symbolic regions** allow discrete-like reasoning to emerge from continuous operations.

1.4 Contributions

This paper provides:

1. A formal topological model of LLM computation.

2. A definition of the preferred direction function and its relation to reasoning.
 3. A continuous analogue of unification grounded in attention.
 4. An attractor-based account of meaning and hallucination.
 5. A unification of symbolic and neural reasoning frameworks.
 6. Architectural implications for explicit reasoning.
-

2 Background and Related Work

This section integrates constraint systems, probabilistic computation, symbolic AI, and modern neural architectures.

2.1 Constraint Satisfaction and Problem Domains

A problem domain can be described as:

- A universe of information (N),
- A constrained subset ($P N$),
- A series of constraints (C_i),
- A solution space that shrinks as constraints accumulate.

LLMs exhibit similar behavior: context narrows semantic possibilities.

2.2 Probabilistic Computation in Brains and Machines

All physical computation involves noise:

- neurons fire stochastically,
- transistors experience thermal noise,
- quantum fluctuations influence all systems.

Brains and LLMs both implement *probabilistic computation refined by constraints*.

Thus, “LLMs can’t reason because they are probabilistic” is unsound.

2.3 Transformers and Attention

Attention is a continuous pattern-matching mechanism:

- Queries select which Keys matter,
- Values transmit semantic information,
- Multi-head attention performs parallel soft constraints.

This parallels symbolic unification in spirit, not mechanism.

2.4 Markov Blankets and Emergent Boundaries

Markov blankets define statistical boundaries in probabilistic graphical models. In LLMs, clustering in embedding space creates regions with high internal coherence and partial independence—pseudo-Markov blankets.

These regions behave symbolically.

2.5 Symbolic Reasoning and Prolog

Prolog unifies patterns via exact symbolic matching. LLMs unify via similarity in embedding space. Both systems:

- match patterns to constraints,
- propagate bindings,
- perform structured inference.

One is discrete, the other continuous.

3 Foundations: The Semantic Manifold and Directional Flow

We now introduce the formal mathematical grounding.

3.1 Semantic Manifold

Let:

- (M^d) : the set of valid activations of the model.

We interpret:

- each point $(x \in M)$ as a **semantic state**,
 - trajectories (x_0, x_1, \dots, x_T) as **reasoning paths**.
-

3.2 Constraint Sets

Context (c) induces a constraint set: $[(c) M]$ representing all states consistent with context.

Strongly constrained tasks produce narrow $((c))$.

3.3 Preferred Direction Function

Definition 1.

A **preferred direction function** is: [$D : M \rightarrow T(M)$,] where ($T(M)$) is the tangent bundle of (M).

The model performs: [$x_{t+1} = x_t + t D(x_t, c_t)$.]

This is the *core* of LLM reasoning.

Attention **implements** this function.

3.4 Diagram: Model Overview



Diagram 0

4 Soft Unification and the Transformer as a Dynamical System

4.1 Soft Unification Operator

Attention computes: [$U_{\{ \}}(q) = _i _i(q, k_i) v_i$,] with [$_i = .$]

This is continuous unification:

- symbolic → discrete substitution,
 - soft → similarity-weighted synthesis.
-

4.2 Multi-Head Attention as Constraint Composition

Each head captures:

- relations,
- analogies,
- structural patterns.

Stacked layers → deeper constraints.

4.3 Directional Dynamics

Attention biases the direction of semantic flow: [$D(x,c) = (x,c)$.]

Thus the Transformer is a **dynamical system**.

5 Proto-Symbols, Attractors, and Statistical Boundaries

5.1 Proto-Symbol Regions

Clusters in embedding space form regions where:

- trajectories enter and remain stable,
- outputs are semantically narrow,
- internal correlations are strong.

Definition 2.

A region (R_M) is a **proto-symbol** if:

- it is an attractor-like set,
 - outputs correspond to consistent semantic categories.
-

5.2 Diagram: Attractor Structure




Diagram 1

5.3 Markov-Blanket-Like Boundaries

Inside (R_-):

- states predict each other well,
- external influence is mediated through boundary states.

This gives symbolic-like integrity.

6 Symbolic Systems and Continuous Systems: A Unified Unification Theory

6.1 Prolog's Discrete Unification

[(t_1, t_2) .]

Search = tree exploration.

6.2 Transformer's Continuous Unification

[U_{\{\}} : ^d \wedge ^d.]

Inference = manifold flow.

6.3 Correspondence

Prolog	Transformer
Terms	Embeddings
Unification	Soft similarity matching
Backtracking	Parallel weighted evaluation
Substitutions	Continuous blending
Discrete search	Dynamical system trajectory

Both instantiate the *same abstract pattern-matching operation* in different computational substrates.

7 Architectural Implications for Explicit Reasoning

7.1 Directional Modulation

Modify: [D_{\{\}}(x,c) = D_{\{\}} + k D^{\{\}}(k).]

Sources of (D_{\{\}}):

- symbolic validators,
- type systems,
- tool outputs,

- RL reward shaping,
 - retrieval grounding.
-

7.2 Topology-Aware Regularization

Encourage:

- clean attractors,
- clear boundaries,
- valid trajectories.

Penalize:

- drift into weak regions,
 - wandering away from grounding.
-

7.3 Hybrid Neuro-Symbolic Systems

Symbolic components define tighter constraint sets: [_{}(c).]

These reshape the manifold traversal.

8 Conclusion

We have shown that LLM reasoning emerges from:

- semantic manifolds,
- constraint-induced subspaces,
- attention-mediated direction fields,
- proto-symbolic attractors.

This framework unifies symbolic and neural accounts and provides a principled basis for building explicit reasoning modules.

Appendix A: Deeper Background on Topology, Computation, and Reasoning

This appendix contains deeper discussions that exceed the main paper's constraints.

A.1 The Fundamental Computation: Direction over a Topology

Given a topology, the only primitive computation is: [D(x,c),] the preferred direction.

Everything else—optimization, inference, reasoning—is emergent trajectory behavior.

A.2 Attractors as Meaning Structures

Semantic stability = attractor. Symbolic meaning = stable manifold.

A.3 Hallucinations as Chaotic Excursions

Hallucinations occur when trajectories enter underconstrained basin regions.

A.4 Category-Theoretic Interpretation

- Embeddings = objects in an enriched category
 - Attention = enriched morphisms
 - Reasoning = sequence of morphisms forming a path
 - Proto-symbol regions = limits or terminal objects
-

References

(As in the original document, plus new citations as needed.)